

LONDON CHAPTER · ENTREPRENEUR FIRST OFFICE

Empathetic Agents

Hackathon Handbook

Thursday, 30 July 2026 · 17:00 – 22:00

Entrepreneur First office, London

Two tracks · free AI credits · food & drinks

“You’ve built an agent — but have you built one that knows when to stay quiet?”



EverSettled × ElevenLabs × Lovable

This handbook is also available as a Markdown file (Empathetic-Agents_Handbook.md), formatted for easy processing by AI assistants.

1. Welcome

You've built an agent — but have you built one that knows when to stay quiet? Or one that catches itself before it accidentally gives legal advice?

Welcome to **Empathetic Agents**, an evening hackathon about building AI that shows up when life gets hard: a diagnosis, a loss, a layoff, a night you can't sleep. Most software treats a person in crisis like a support ticket. Tonight we want to see what it takes to build something better — and safer.

Everything you need is in this handbook: the two tracks, your AI key, your ElevenLabs and Lovable credits, how to submit, and how you'll be judged.

2. Schedule

Approximate and subject to change — check the night before.

TIME	WHAT
5:00–5:30 PM	Doors open, food, team formation
5:30–5:45 PM	Welcome and challenge briefing
5:45–8:45 PM	Hacking
8:45–9:15 PM	Judging
9:15–9:45 PM	Demos
9:45–10:00 PM	Awards and close

Three hours of building. Best teams from both tracks will be invited to the stage to demo and share their approach.

3. The tracks

Pick a track. Both leave room to be creative — one just has a wider brief than the other.

Track 1 — Presence

An open, exploratory track. Build an agent that helps someone through something hard: not by answering one question well, but by **being there for the stretch of it**. Two directions to take it, or blend them:

- **A copilot** — an agent that helps someone find distraction, meaning, or joy alongside the hard thing they're going through. A bucket list generator for someone with a new diagnosis, a late-night chat companion, a way to plan a good day in the middle of a bad month.
- **The infrastructure to track progress and make decisions** — the system underneath: what does it remember, and how does it decide when to check in versus when to leave someone alone?

Judged on **originality**, **technical execution**, and **how genuinely the agent responds** to a person going through something hard. There's a template to get you started, but get creative with the direction.

Track 2 — Guardrails

A narrower brief, same freedom to solve it your way. Build the layer that **catches an empathetic agent when it gets something wrong**.

- **AI guardrails for dealing with death** — the checks that stop an agent overstepping, sounding falsely certain, or saying the wrong thing around death, grief, legal, or financial matters.
- **An outlier detector** — a layer that spots when someone's signal is a real crisis, not normal frustration, and routes it differently.

Judged on **technical rigor**, **how well it catches the failure case**, and **how it holds up against inputs designed to break it**.

4. How you are judged

The same four criteria apply to both tracks, and all four are **weighted equally**:

- **Originality** — how fresh or unexpected is the idea and approach.
- **Technical Execution** — how far did they get, and how well does it actually work.
- **Applicability** — would this genuinely help someone, or matter to EverSettled specifically.
- **Fit to the Brief** — judged live, off what you see in the demo, no test cases needed.

For Presence teams: does it actually feel like it understands the person, not just the prompt.

For Guardrails teams: does it actually catch the bad behaviour it's meant to catch.

5. A note on care

You're designing for people at their most vulnerable. The strongest entries are honest about their limits: they don't pretend to be a therapist or a doctor, they recognise signs of real crisis and signpost proper help, and they hand off to humans when that's the kind thing to do. We reward care and responsibility as much as cleverness.

6. Your AI key

Your team has one API key with a **\$50 budget** for the event. It works with any OpenAI SDK or tool — you only change two things: the **API key** and the **base URL**.

Base URL: `https://4.231.223.10.nip.io`

API key: `sk-...` (the one your team was given – don't share or commit it!)

Python

```
from openai import OpenAI
```

```
client = OpenAI(
    api_key="sk-...your-team-key...",
    base_url="https://4.231.223.10.nip.io",
)
r = client.chat.completions.create(
    model="gpt-5.4-mini",
    messages=[{"role": "user", "content": "Give me three hackathon project ideas."}],
)
print(r.choices[0].message.content)
```

JavaScript / TypeScript

```
import OpenAI from "openai";
```

```
const client = new OpenAI({
  apiKey: "sk-...your-team-key...",
  baseURL: "https://4.231.223.10.nip.io",
});
const r = await client.chat.completions.create({
```

```

model: "gpt-5.4-mini",
messages: [{ role: "user", content: "Hi!" }],
});

```

Anything that speaks the OpenAI API works the same way (LangChain, Vercel AI SDK, curl, ...).

Available models

MODEL	USE IT FOR	API
gpt-5.4-mini	Start here — fast, cheap, strong	Chat Completions
gpt-5.4-nano	High-volume / simple calls	Chat Completions
gpt-5.4	Harder tasks	Chat Completions
gpt-5.5	Hardest reasoning (pricier — use deliberately)	Chat Completions
gpt-5.6-terra	Newest generation, small tier	Chat Completions
gpt-5.6-luna	Newest generation, mid tier	Chat Completions
gpt-5.6-sol	Newest gen, flagship (priciest — use deliberately)	Chat Completions
gpt-5.3-codex	Coding agents	Responses API only (client.responses.create)
gpt-audio	Speech in and/or speech out	Chat Completions, modalities=["text", "audio"]
gpt-realtime-2.1	Live voice conversations	WebSocket: wss://4.231.223.10.nip.io/v1/realtime?mode=l=gpt-realtime-2.1
gpt-realtime-2.1-mini	Live voice, cheaper	Same, with model=gpt-realtime-2.1-mini

Audio example — get a spoken answer back

```

import base64

r = client.chat.completions.create(
    model="gpt-audio",
    modalities=["text", "audio"],
    audio={"voice": "alloy", "format": "wav"},
    messages=[{"role": "user", "content": "Welcome our demo audience in one sentence."}],
)
open("answer.wav", "wb").write(base64.b64decode(r.choices[0].message.audio.data))

(Send audio in the same way with a {"type": "input_audio", ...} content part.)

```

Your budget

- Your key has a **hard \$50 cap** — when it's used up, requests are rejected. Nothing you can break, but plan for it: a runaway agent loop on gpt-5.6-sol or gpt-5.5, or long realtime sessions, are the fastest ways to burn it. Develop on gpt-5.4-mini, switch up only where it matters.
- If a request fails with a budget error and you believe something's wrong, come find the organizers.

7. ElevenLabs credits

Every participant gets **1 month of the ElevenLabs Creator plan** (worth \$22) — ideal for voice-first agents, natural speech, and anything where tone carries the empathy.

Get your key from us at the event, or email tobias.budig@eversettled.com.

8. Lovable credits

You also get free **Lovable Pro Plan 1 (100 credits)** for building front-ends fast.

How it works

- The code gives access to Pro Plan 1 (100 credits) at no cost.
- Unlocks all Pro Plan 1 features for the duration of the plan.
- The code must be redeemed by the end of the event, and should not be shared with anyone outside of the event.

How to apply the code

- 1 Go to lovable.dev.
- 2 If you don't have an account, click "Get started" → Create an account.
- 3 If you have an account, go to Settings → Plans & Credits.
- 4 Select **Pro Plan 1 (100 credits)**. Make sure to choose the monthly plan.
- 5 At checkout, enter your discount code: **COMM-HACK-3YGM**

If you already have a paid account, just create a new workspace to redeem the code.

Please note: you may be asked to enter payment details, but you have one month to cancel, and can cancel as soon as the account is activated.

9. Wifi

Network: Shoreditch Exchange - Member
Password: ShoreditchX!

10. Submitting your project

One entry per team via the **submission form**. You'll need:

- **Team name**
- **Team members**
- **GitHub URL**
- **Presentation URL**

The **submission form link will be shared in the WhatsApp group**, along with all other details and updates during the event — so make sure you've joined.

Submit before judging begins at **8:45 PM**. Best teams from both tracks are invited to the stage to demo and share their approach at 9:15 PM.

11. House rules

- Be respectful and inclusive — all backgrounds and skill levels are welcome.
- **Clean up after yourself.** We're guests in this space: bin your rubbish, clear your table, and leave it as you found it.
- Build during the event; using open-source tools and starter code is encouraged.
- **Don't use real personal or sensitive data** in your prototypes — use your own samples or synthetic data.

- Redeem the AI, ElevenLabs, and Lovable credits for this event only, and don't share codes outside it.
- Respect open-source licences.

Our partners



EverSettled is an AI-native estate-settlement platform helping families settle the affairs of a loved one after they die.

Today, estate settlement takes 18 months, costs tens of thousands of dollars, and requires up to 900 hours of unpaid administrative work — often forcing executors to absorb out-of-pocket expenses while grieving. That’s exactly the kind of hard moment tonight is about.

ElevenLabs — voice AI, powering the way these agents sound. elevenlabs.io

Lovable — build front-ends fast. lovable.dev

Empathetic Agents · London chapter · EverSettled × ElevenLabs × Lovable